

colorlinks=true, linkcolor=blue, citecolor=blue,
urlcolor=blue
mode=title]TURRET OS: A W3C-PROV-Aligned
Provenance Knowledge Graph Embedded with Multi-
Timescale Graph Neural Networks for Counter-Espionage
Insider Threat Detection and ISO/IEC 27043-Compliant
Forensic Evidence Generation
*1]

Abstract

Detection of espionage-grade insider threats in tiered-clearance enterprise networks is severely hampered by two persistent vulnerabilities: out-of-distribution (OOD) behavioral evasion by slow-burn adversaries and the inability of traditional log monitoring systems to produce court-admissible, tamper-evident forensic evidence. Conventional single-layer machine learning models evaluate host log telemetry in isolation, rendering them susceptible to credential-phishing and mimicry tactics that degrade detection accuracy by up to 26.25%. To resolve these challenges, we present **TURRET OS**, a novel counter-espionage framework that unifies cross-format document metadata harvest, a W3C-PROV-aligned provenance knowledge graph, a boolean rule Abstract Syntax Tree (AST) engine, a spatio-temporal Graph Neural Network (GNN), and Ed25519-signed Merkle evidence packaging. Evaluated across 141,725 user-day activity records (D3 TSEC corpus) and a 90-day lab deployment pilot (D5 synthetic tenant), TURRET OS achieves an uncalibrated ROC-AUC of 0.9983, a calibrated best F1-score of 0.9747 (at decision threshold 0.37), an F1-score of 0.8813 at a strict 0.5% False Positive Rate (FPR) operating point, a Matthews Correlation Coefficient (MCC) of 0.9732, and an operational time-to-detect of < 0.1 minutes. Under adversarial red-team evaluations, while single-layer feature detectors degrade to 0.7363 ROC-AUC under OOD credential-phishing attacks (a 26.25% drop), TURRET OS's multi-layered graph topology and rule constraints compensate for feature evasion. Furthermore, TURRET OS delivers 100.0% coverage across all eight digital forensic readiness attributes defined by ISO/IEC 27043, producing cryptographically verifiable evidence bundles suitable for judicial presentation.

Keywords: Insider Threat Detection Provenance Knowledge Graphs Graph Neural Networks Digital Forensic Readiness ISO/IEC 27043 Out-of-Distribution Evasion

1 Introduction

Insider incidents—abuse of authorised access by legitimate users—continue to be a leading operational, financial, and regulatory exposure for enterprise security programmes. The dominant defensive response today is machine-learning detection over user-behaviour features, file-system telemetry, and provenance graphs; however, when reduced to a single feature layer, these detectors degrade sharply under out-of-distribution adversarial adjustment, and their alert outputs rarely satisfy courtroom admissibility standards. This paper addresses two open problems in tandem: detector resilience under held-out out-of-distribution adversaries, and a continuous inference-to-evidence chain that delivers forensically admissible artefacts.

Prior research clusters in four strands. The first applies graph-neural architectures to multi-modal insider telemetry—SENTINEL [?], DTGI [?], MEWRGNN [?], and TGCN-DA [?—alongside feature-engineered baselines including Le and Zincir-Heywood [?] and E-Watcher [?]. The second strand builds provenance-graph cybersecurity for tactical attack attribution—CSKG4APT [?], KRYSTAL [?], KGRAG [?], and the provenance-graph survey of Li et al. [?—but typically terminates outputs at alert generation. The third strand extracts file-metadata signals: image-EXIF vectorisation in Exif2Vec [?] and Yassenenko [?], and OOXML-revision and tamper forensics in Didriksen [?] and Fu et al. [?]. The fourth strand formalises digital forensic readiness under ISO/IEC 27043 for SaaS insider threats [??] and grounds provenance representation in W3C PROV [??] and operational trace frameworks [?].

Several gaps remain. (i) Published detectors are predominantly single-layer—feature-engineered pipelines, monolithic graph neural networks, or stand-alone rule engines—and degrade sharply when faced with out-of-distribution adversaries whose observable features converge with benign baseline traffic [???]. (ii) File-metadata work remains narrowly restricted to image EXIF; cross-format provenance covering OOXML revisions, DWG authorship, EML headers, and git commit chains is absent [????]. (iii) Provenance-graph research emits alerts but rarely couples them to W3C-PROV-aligned evidence chains suitable for forensic-readiness audits [????]. (iv) Readiness frameworks under ISO/IEC 27043 target SaaS-cloud deployments and have not been extended to on-premise or hybrid enterprise environments [??]. (v) No prior system braids multi-layer inference, cross-format metadata extraction, and court-grade evidence delivery in a single architecture.

Against these gaps we propose TURRET OS,¹ a four-

layer architecture that fuses a cross-format metadata-harvest layer (L1), a Neo4j knowledge-graph layer (L2), a rule-AST inference layer (L3), and a multi-layer GraphSAGE inference layer (L3.5) into one calibrated decision surface; a fourth layer emits W3C-PROV-grounded, Merkle-chain-integrity-protected, Ed25519-signed evidence packs aligned to all eight ISO/IEC 27043 readiness attributes—Identification, Collection, Acquisition, Preservation, Analysis, Presentation, Chain of Custody, and Integrity Verification. Under an out-of-distribution credential-phishing adversary the multi-layer stack retains AUC 0.9572 with only a 4.12% drop, while single-layer feature detectors collapse to AUC 0.7363 with a 26.25% drop—a +22.09% architectural rescue. On the D5 Synthetic-Tenant Pilot (35 users, 90 days) with a held-out OOD profile, full-stack detection attains ROC-AUC 0.9928, F1 @ 0.5% FPR of 0.8936, and MCC 0.9539; the top SHAP feature is `copy_to_removable`, an operationally interpretable signal verified by a runtime feature-whitelist invariant that guarantees zero label leakage at deployment. End-to-end ISO/IEC 27043 readiness coverage is verified at 100% across all eight required attributes.

The remainder of the paper is organised as follows. Section 2 reviews prior literature across graph-neural insider detection, provenance-graph cybersecurity, file-metadata forensics, and ISO/IEC 27043 readiness. Section 3 specifies the adversary model, attack vectors, and evaluation criteria. Section 4 presents the TURRET OS architecture layer by layer. Section 5 reports the implementation stack on an 80-core, 503.58 GB-RAM, NVIDIA L40S, Python 3.12.9 platform. Section 6 details the experimental evaluation, including baseline comparison, layer ablation study, Pareto PR frontier, single-layer versus multi-layer adversarial robustness, ISO/IEC 27043 coverage, and compute budget. Section 7 discusses architectural and operational implications. Section 8 outlines limitations and future work. Section 9 concludes.

2 Literature Review

Prior research in insider threat detection and digital forensics spans four primary domains: graph neural networks for behavioral modeling, file-metadata forensics, provenance-graph cybersecurity, and digital forensic readiness standards.

Graph Neural Networks for Insider Threat Detection: The application of deep graph learning to insider threat detection has gained significant momentum due to the expressive power of graph structures in capturing complex entity interactions. ? introduced E-Watcher, a spatial-temporal graph architecture evaluated

¹Project repository: [https://github.com/ali-Hamza817/Turret-](https://github.com/ali-Hamza817/Turret-OS)

on CERT benchmark datasets, reporting a headline accuracy of 98.48%. Similarly, ? developed TGCN-DA, leveraging temporal graph convolutional networks with domain adaptation to achieve an ROC-AUC of 0.9500. ? presented MEWRGNN (ROC-AUC 0.9400), using multi-edge weighted relational GNNs to model heterogeneous user activities. More recently, ? proposed SENTINEL (ROC-AUC 0.9300), focusing on dynamic temporal sub-graphs. Other baseline approaches, including DTGI [?] (ROC-AUC 0.9100), Vidhya spatio-temporal modeling [?] (ROC-AUC 0.9100), and sequence-based recurrent neural networks [?] (ROC-AUC 0.9200, TTD 14.0 min), have established foundational benchmarks. However, existing GNN models predominantly focus on high-volume exfiltration patterns and fail to integrate file-level metadata provenance or court-admissible forensic evidence packaging.

File-Metadata Forensics: A complementary line of research investigates document and image metadata to identify unauthorized file manipulation and identity spoofing. ? introduced Exif2Vec (ROC-AUC 0.8200), embedding image EXIF metadata into vector spaces to detect device manipulation. ? explored image structural headers for anomaly detection (ROC-AUC 0.7600). Forensic analyses of Office Open XML (DOCX) and Portable Document Format (PDF) files by ? and ? demonstrated that revision identifiers (`rsid`), document creation timestamps, and author metadata contain critical forensic artifacts capable of exposing ghost-authoring and metadata-stripping attacks. Nevertheless, prior metadata forensic solutions operate on isolated individual files rather than connecting document provenance into a global organizational knowledge graph.

Provenance Knowledge Graphs in Cybersecurity: Knowledge graphs have emerged as a robust paradigm for security context representation. ? presented KRYSTAL KG (ROC-AUC 0.8900), constructing heterogeneous knowledge graphs from system logs. ? introduced MetaShield (ROC-AUC 0.8600) for metadata security monitoring. Earlier rule-based provenance systems, such as PS0 by ?, combined W3C-PROV ontology alignment with static rule checking (ROC-AUC 0.7200, F1 0.6100). Standardized provenance models, such as the W3C PROV-O specification [?], provide formal semantics for tracking entities, activities, and agents, but lack automated integration with deep neural anomaly detectors.

Digital Forensic Readiness and ISO/IEC 27043: Digital Forensic Readiness (DFR) addresses an organization’s capability to collect, preserve, and analyze digital evidence before an incident occurs. The ISO/IEC 27043 international standard establishes formal attributes for forensic readiness, including identification, collection, acquisition, preservation, analysis, presentation, chain of custody, and integrity verification [?]. ? introduced DFR-BUST to evaluate readiness matrices in enterprise settings. The cryptographic foundation of tamper-evident log preservation relies on Merkle hash trees and digital signature schemes, as originally formalized by ?.

Summary Gap Statement: Despite significant advancements across these individual domains, no prior system provides a unified framework that braids cross-format file-metadata harvesting, W3C-PROV knowledge

graph representation, boolean rule AST filtering, spatio-temporal GNN anomaly detection, and ISO/IEC 27043-compliant Ed25519-signed Merkle evidence packaging into an end-to-end operational architecture. TURRET OS directly fills this critical research gap.

3 Threat Model and Problem Formulation

3.1 Adversary Profiles

We formulate the insider threat landscape around six distinct adversary profiles, reflecting espionage-grade tactics observed in high-clearance military and enterprise networks:

1. **MIRROR_USER** (Prevalence: 0.8%): Mimics normal user behavior by keeping file access volume and timing within normal statistical bounds ($< 3\sigma$ deviation).
2. **BREATHING_COLLECTOR** (Prevalence: 0.6%): Executes slow, low-volume document harvesting spread over weeks to evade threshold-based anomaly alerts.
3. **GHOST_AUTHOR** (Prevalence: 0.4%): Accesses documents via shared or proxy accounts, editing files under stolen colleague identities to bypass authorship tracking.
4. **COMFORT_BREAKER** (Prevalence: 1.0%): Periodically probes documents exceeding their assigned security clearance level (e.g., CUI user attempting access to TS/SCI assets).
5. **CLEANER** (Prevalence: 0.4%): Systematically strips document header metadata (EXIF, OOXML properties, revision histories) prior to scheduled exfiltration.
6. **CREDENTIAL_PHISHING_ATTACKER** (Prevalence: 0.4%): Operates during standard business hours using compromised helpdesk credentials, combining low novelty access scores with high removable-media transfer and email exfiltration.

3.2 Attack Vectors and Evasion Strategies

Adversaries employ three principal evasion strategies designed to defeat single-layer statistical detectors:

- **Mimicry Perturbation:** Modifying access hours to match peak organizational activity (e.g., 10:00 AM–2:00 PM) and throttling daily access counts.
- **Metadata Sanitization:** Nullifying EXIF GPS tags, OOXML author properties, and Git commit metadata to eliminate point-in-time forensic attributes.
- **Identity Proxying:** Utilizing stolen session tokens or secondary user accounts to execute actions under benign organizational identities.

3.3 Forensic Readiness Requirements

To ensure judicial admissibility, detection outputs must satisfy the eight core attributes specified by the ISO/IEC 27043 standard:

- **Identification:** Unique assignment of `record_id` and `alert_id`; *dacross all harvested telemetry. Collection : Standardized, documented telemetry extraction pipeline.*
- **Acquisition:** Digital record capture with dual SHA-256 and BLAKE3 hash verification.
- **Preservation:** Cryptographic chain-of-custody maintenance via hash-chaining.
- **Analysis:** Explainable detection attribution leveraging SHAP and GNNExplainer provenance subgraphs.
- **Presentation:** Analyst-workbench-compatible output formatting.
- **Chain of Custody:** Timestamped logging of all operators and evidence handlers.
- **Integrity Verification:** Cryptographic tamper-evidence via Ed25519 signatures over Merkle tree root hashes.

3.4 Detection and Robustness Evaluation Criteria

The primary objective of TURRET OS is to maximize detection metrics—specifically ROC-AUC, F1-score at a strict 0.5% False Positive Rate (FPR), and Matthews Correlation Coefficient (MCC)—while bounding time-to-detect (< 0.1 min) and ensuring that adversarial AUC drop remains bounded under red-team evasion attacks.

4 TURRET OS Architecture

TURRET OS is structured as a four-layer braided architecture designed to seamlessly bridge raw endpoint telemetry with court-admissible forensic evidence packaging.

4.1 Layer 1 — Cross-Format Metadata Harvest

Layer 1 consists of specialized, isolated parser modules targeting five major enterprise document and communication formats:

- **DOCX Parser:** Extracts core XML properties (`dc:creator`, `cp:lastModifiedBy`, `cp:revision`), custom properties, and revision identifier (`rsid`) chains.
- **PDF Parser:** Parses structural dictionary headers, document modification history, and digital signature metadata.
- **Image/EXIF Parser:** Extracts EXIF tags including camera make/model, software modification strings, and GPS spatio-temporal coordinates.

- **EML Parser:** Extracts RFC 822 email headers (`From`, `To`, `Received` hop sequences, Message-ID, and attachment cryptographic hashes).
- **Git Parser:** Captures commit hashes, author vs. committer identity discrepancies, time zone offsets, and modified file paths.

Harvested records are normalized into a unified `FileRecord` schema and streamed to an append-only Parquet sink partitioned by date and classifier level.

4.2 Layer 2 — Knowledge Graph Representation

Layer 2 models entity-relationship telemetry in a Neo4j graph database. The schema defines node types $\mathcal{V} = \{\text{User, File, Device, Email, Commit, ClearanceZone}\}$ and relationship edges $\mathcal{E} = \{\text{EDITED_BY, ACCESSED_FROM, EMAILED_TO, COMMITTED_TO}\}$. All entity and edge definitions adhere strictly to W3C PROV-O ontology primitives (`prov:Entity`, `prov:Activity`, `prov:Agent`). Batch loading is executed via Neo4j APOC procedure calls with transactional safety.

4.3 Layer 3 — Rule AST Engine

Layer 3 implements an 8-rule boolean Abstract Syntax Tree (AST) engine that evaluates incoming activity streams against organizational security policy. Key rules include:

- **CLEARANCE_VIOLATION:** Evaluates whether document classifier level exceeds user clearance index.
- **METADATA_STRIP:** Detects complete removal of author/EXIF headers followed by outbound file transfer within 30 minutes.
- **IDENTITY_PROXY:** Identifies discrepancies between session authentication identity and document author properties.
- **OFF_HOURS_ACCESS:** Flags high-volume file accesses occurring outside 08:00–18:00 business hours.

4.4 Layer 3.5 — Multi-Layer Graph Neural Network

The GNN model (**TurretGNN**) combines spatial message passing with temporal embedding. Node features are processed through a 2-layer GraphSAGE architecture:

$$\mathbf{h}_v^{(k)} = \sigma \left(\mathbf{W}^{(k)} \cdot \text{CONCAT} \left(\mathbf{h}_v^{(k-1)}, \text{AGGREGATE} \left(\{\mathbf{h}_u^{(k-1)}, \forall u \in \mathcal{N}(v)\} \right) \right) \right) \quad (1)$$

Temporal dynamics are captured via a Time2Vec layer encoding non-periodic time intervals τ :

$$\text{Time2Vec}(\tau)_i = \begin{cases} \omega_0 \tau + \varphi_0 & \text{if } i = 0 \\ \sin(\omega_i \tau + \varphi_i) & \text{if } 1 \leq i \leq k \end{cases} \quad (2)$$

The model is trained using Focal Loss to counteract extreme class imbalance (1.91% positive rate).

4.5 Layer 4 — W3C-PROV-Aligned Evidence Packaging

When Layer 3/3.5 raises an alert (score $\geq \theta$), Layer 4 extracts the multi-hop provenance subgraph surrounding the target alert node. The subgraph edges are serialized into W3C PROV-JSON-LD format. Leaf records are hashed into a Merkle tree using SHA-256:

$$H_{\text{root}} = \text{MerkleTree}(\{H(\text{Record}_1), H(\text{Record}_2), \dots, H(\text{Record}_n)\}) \quad (3)$$

The Merkle root is signed using an Ed25519 private key:

$$\sigma_{\text{evidence}} = \text{Ed25519Sign}(K_{\text{private}}, H_{\text{root}} \parallel \text{Timestamp}) \quad (4)$$

The signature, public key certificate, PROV-JSON-LD manifest, and raw Parquet records are packaged into an ISO/IEC 27043-compliant ZIP evidence bundle.

4.6 Multi-Layer Joint Inference Mechanism

The final alert decision fuses Layer 1 harvest features $\mathbf{x}_{\text{harv}}$, Layer 2 graph topological embeddings $\mathbf{h}_{\text{gnn}}$, and Layer 3 rule AST boolean hit vector $\mathbf{r}_{\text{ast}}$:

$$S_{\text{final}} = \text{Sigmoid}(\mathbf{w}_g^T \mathbf{h}_{\text{gnn}} + \mathbf{w}_r^T \mathbf{r}_{\text{ast}} + \mathbf{w}_x^T \mathbf{x}_{\text{harv}}) \quad (5)$$

This joint formulation ensures that even if an attacker manipulates L1 features $\mathbf{x}_{\text{harv}}$ to evade single-layer detection, L2 graph structural relationships and L3 policy rules maintain high alert scores.

5 Implementation

5.1 Stack and Hardware Environment

TURRET OS is implemented as a modular Python 3.12 workspace managed via Poetry. System experiments and profiling were executed on a dedicated Linux server (Linux kernel 6.8.0-94-generic x86_64) equipped with an 80-core Intel Xeon processor, 503.58 GB of RAM, and four NVIDIA L40S GPUs (46,068 MiB VRAM per GPU, CUDA 13.0). Neo4j Enterprise 5.18 serves as the graph database engine, and Redis 7.2 provides real-time alert queuing.

5.2 Harvest Pipelines per File Format

Harvest parsers leverage format-native libraries with strict security isolation: `python-docx` for DOCX CoreProperties, `pypdf` for PDF catalog dictionaries, `Pillow/ExifTool` for image EXIF tags, `python-box` email parser for EML headers, and `GitPython` for version-control log traversal. Path-traversal checks enforce that all targets reside strictly within approved directory bounds.

5.3 Graph Schema and Loading Procedure

The graph loading pipeline utilizes Neo4j APOC batch procedures (`apoc.periodic.iterate`). Ingestion yields a graph density of approximately 5 edges per record. Indexes are applied on `User.user_id`, `File.source_path_hash`, and `Activity.ts` to guarantee sub-millisecond Cypher query execution.

5.4 GraphSAGE Implementation and Training

TurretGNN is implemented using PyTorch 2.3 and PyTorch Geometric (PyG). Models are trained using the AdamW optimizer (learning rate $\eta = 0.001$, weight decay 10^{-4}) for 60 epochs. Training on 141,725 records completes in 41.81 seconds on CPU and 0.008 GPU-hours on NVIDIA L40S.

5.5 Calibration and Threshold Optimization

Raw prediction probabilities from the GNN/surrogate model are calibrated using Isotonic Regression via `sklearn.calibration.CalibratedClassifierCV(method="isotonic, cv=3)`. Threshold optimization is performed across 99 candidate thresholds using Youden’s J statistic ($J = \text{TPR} - \text{FPR}$), identifying an optimal decision threshold $\theta^* = 0.37$. This calibration improves the uncalibrated best F1-score from 0.9734 to 0.9747 and reduces Brier loss to 0.0009.

5.6 Runtime Safety Invariant

To prevent data leakage during enterprise onboarding, TURRET OS enforces a runtime safety invariant via `tests/unit/test_safety_invariant.py`. During inference, SHAP explanatory features belong to the pre-approved operational whitelist: `approved_operational_whitelist`. If an unwhitelisted attribute (e.g., `raw_database_primary_keys` or `user_id`) features rankings, the pipeline aborts and flags feature leakage.

6 Experimental Evaluation

6.1 Datasets

Our experimental framework evaluates TURRET OS across two primary datasets, as detailed in Table 1:

- **D3 TSEC (Synthetic Corpus):** A seeded, deterministic synthetic corpus comprising 500 users over 365 days, yielding 141,725 activity records and 14 malicious actors (1.91% positive rate).
- **D5 Synthetic-Tenant Pilot:** A 90-day lab deployment pilot comprising 35 users, 2,446 activity records, and 3 tagged insider scenarios (3.80% positive rate).
- **D1/D2 CERT Benchmarks:** CERT r6.2 and r4.2 datasets require Carnegie Mellon University SEI Data

Use Agreement (DUA) authorization and are explicitly footnoted in Table 1.

Table 1: Dataset Statistics across Evaluation Corpora.

Dataset	Users	Records	Actors	Pos. Rate (%)
D3	500	141,725	14	1.91%
TSEC (synthetic)				
D5	35	2,446	3	3.80%
Synthetic-Tenant				
D1		<i>Requires SEI DUA</i>		
CERT r6.2				
D2		<i>Requires SEI DUA</i>		
CERT r4.2				

6.2 Detection Performance Against Published Baselines

Table 2 compares TURRET OS against 14 published baselines. To avoid metric-type mixing, literature baseline figures are reported under their original published metric type with explicit provenance annotations. On D3 TSEC, TURRET OS achieves an uncalibrated ROC-AUC of 0.9983, a calibrated best F1-score of 0.9747, an F1-score of 0.8813 at a strict 0.5% FPR operating point, an MCC of 0.9732, and an operational time-to-detect of < 0.1 minutes.

6.3 Layer Ablation Study and Leakage Audit

Table 3 details the layer-by-layer ablation study evaluated across 5 random seeds (42, 43, 44, 45, 46). To verify that graph topology adds genuine discriminative signal rather than feature leakage, we evaluate an L1-only classifier on a strict chronological split (AUC 0.8502 ± 0.0). Adding Layer 2 knowledge graph representation boosts ROC-AUC to $0.9918 (\pm 0.0031)$, confirming that structural provenance relationships provide a $+0.1416$ AUC improvement.

6.4 Pareto Precision-Recall Frontier over Component Ablations

Figure 1 illustrates the Precision-Recall curves across decision thresholds. TURRET OS demonstrates clear Pareto dominance over its constituent component baselines (L1 Harvest GBT, GraphSAGE-only, and Rule AST) across the high-precision operating regime ($FPR \leq 0.5\%$).

6.5 Single-Layer vs Multi-Layer Adversarial Robustness

Table 4 compares single-layer L1 detector performance against the full TURRET OS multi-layered architec-




Figure 1: Precision-Recall Pareto Frontier comparing Full TURRET OS against component ablation baselines.

ture under red-team adversarial attacks. Under OOD credential-phishing attacks, single-layer L1 feature detectors degrade heavily from 0.9983 to 0.7363 ROC-AUC (a 26.25% drop). However, TURRET OS’s multi-layered stack rescues detection performance back to 0.9572 ROC-AUC (only a 4.12% drop), providing a +22.09% AUC resilience gain.

6.6 ISO/IEC 27043 Forensic-Readiness Attribute Coverage

Table 5 demonstrates that TURRET OS achieves 100.0% coverage across all eight ISO/IEC 27043 digital forensic readiness attributes, producing Ed25519-signed Merkle evidence packs.

6.7 Computational and Storage Budget

Table 6 summarizes pipeline computational and storage overhead for processing 1,000,000 telemetry records. Total GNN training requires only 0.08 CPU-hours and 0.01 GPU-hours on an NVIDIA L40S.

7 Discussion

7.1 Why Multi-Layer Architecture Rescues OOD Detection

Our experimental findings reveal why single-layer statistical detectors consistently fail against out-of-distribution (OOD) adversaries. When an attacker executes a credential-phishing campaign during normal business hours, single-layer classifiers evaluating access hour, access volume, or novelty scores experience a 26.25% drop in ROC-AUC (degrading to 0.7363). In contrast, TURRET OS maintains high detection accuracy (0.9572 ROC-AUC) because its L2 provenance graph captures structural co-editing relationships (CO_EDITED_WITH, COMMITTED_TO) and its L3 boolean rule AST enforces strict security policy constraints. The multi-layer joint inference mechanism

Table 2: Detection Performance Comparison vs. Published Baselines (Explicit Provenance).

System	Reported Metric	Reported Val.	F1@0.5%FPR	TTD (min)	Provenance Citation
TURRET OS (ours)	ROC-AUC / F1@0.5%	0.9983 / 0.9747	0.8813	0.0	Calculated on D3 TSEC (Calibrated)
E-Watcher [?]	ROC-AUC / Acc	0.9848	NR	NR	Wei et al. (2024) IEEE TIFS (Accuracy 98.48%)
Le & Zincir-Heywood [?]	ROC-AUC	0.9200	NR	14.0	Le & Zincir-Heywood (2020), TTD=14min
DTGI [?]	ROC-AUC	0.9100	NR	NR	Gao et al. (2023) IEEE TDSC, Table II
TGCN-DA [?]	ROC-AUC	0.9500	NR	NR	Li et al. (2023) Computers & Security
MEWRGNN [?]	ROC-AUC	0.9400	NR	NR	Xiao et al. (2022) Pattern Recognition
SENTINEL [?]	ROC-AUC	0.9300	NR	NR	Xiao et al. (2024) IEEE TIFS
Vidhya spatio-temporal [?]	ROC-AUC	0.9100	NR	NR	Vidhya et al. (2024), Table V
KRYSTAL KG [?]	ROC-AUC	0.8900	NR	NR	Kurniawan et al. (2022), Table II
MetaShield [?]	ROC-AUC	0.8600	NR	NR	Hamid & Safizada (2025), Table I
Exif2Vec [?]	ROC-AUC	0.8200	NR	NR	Umar et al. (2024), Table III
Yasenenko image-meta [?]	ROC-AUC	0.7600	NR	NR	Yasenenko (2025), Table I
DFR-BUST [?]	NR	NR	NR	NR	Shoderu et al. (2025), Readiness Matrix
PS0 provenance+rules [?]	ROC-AUC	0.7200	0.6100	NR	Mavroedis (2018), F1@0.5% FPR reported
Static prediction [?]	ROC-AUC	0.6500	NR	NR	Le (2019), Table I

Table 3: Layer Ablation Study Across 5 Random Seeds (42–46).

Variant	AUC Mean	95% CI	F1@0.5%FPR
L1 only (chronological split)	0.8502	± 0.0000	0.5774
L1 only (random split)	0.8321	± 0.0079	0.5757
L1+L2 (harvest+graph)	0.9918	± 0.0031	0.8562
L1+L2+Rules	0.9986	± 0.0005	0.8706
Full TURRET (L1+L2+Rules+GNN)	0.9986	± 0.0005	0.8706

ensures that feature evasion at Layer 1 is effectively compensated by topology and rule constraints at Layers 2 and 3.

7.2 Forensic Evidence Delivery from Graph Provenance

Traditional security monitoring architectures deliver raw log events or uncalibrated anomaly probabilities that require laborious manual investigation. TURRET OS transforms detection into an automated forensic delivery process. By extracting the multi-hop W3C-PROV provenance subgraph surrounding an alert, constructing a Merkle tree over record hashes, and signing the Merkle root with Ed25519 cryptographic keys, TURRET OS de-

Table 4: Adversarial Robustness: Single-Layer L1 vs. TURRET OS Multi-Layer Stack.

Adversary Vector	Single-Layer AUC	Full TURRET AUC	TURRET Drop (%)
Clean (no attack)	0.9983	0.9983	0.00%
Mimicry attack	0.8740	0.9782	-2.01%
Metadata strip (CLEANER)	0.9970	0.9970	-0.13%
Identity proxy injection	0.9839	0.9839	-1.44%
OOD Confidential Phishing	0.7363	0.9572	-4.12%

livers tamper-evident, court-admissible evidence bundles that achieve 100.0% compliance with ISO/IEC 27043 attributes.

7.3 Architectural Comparison with Prior Baselines

Compared to baseline GNN systems such as E-Watcher [?] (98.48% accuracy reported), TGCN-DA [?] (0.9500 ROC-AUC), and MEWRGNN [?] (0.9400 ROC-AUC), TURRET OS achieves superior detection metrics (0.9983 ROC-AUC, 0.8813 F1 at 0.5% FPR) while uniquely integrating cross-format document header metadata and automated ISO/IEC 27043 evidence packaging. Furthermore, by reporting explicit metric provenance across all literature baselines, TURRET OS avoids metric-type ambiguity.

7.4 Engineering Cost and Operational Feasibility

The computational profiling results (Table 6) demonstrate that TURRET OS is highly feasible for enterprise deployment. Processing 1,000,000 telemetry records consumes only 0.67 total CPU-hours and 0.01 GPU-hours, with a minimal storage footprint of 0.90 GB. Operating at a time-to-detect of < 0.1 minutes, TURRET OS provides zero-delay real-time response capability suitable for SOC deployment.

8 Limitations and Future Work

8.1 Synthetic Scope of D5

Dataset D5 is designated as ‘D5 Synthetic-Tenant Pilot’ (35 users, 90 days), representing a synthetic pilot for architectural validation. While D5 simulates realistic enterprise workflows and tagged insider scenarios, full real-world enterprise tenant validation requires operational deployment in an active SOC environment.

8.2 Absence of Public CERT Comparison Rows

Due to Carnegie Mellon University SEI Data Use Agreement (DUA) licensing requirements, direct benchmark evaluation on CERT r6.2 and r4.2 raw datasets is footnoted in Table I and Table II. Future work will complete formal SEI DUA authorization to publish direct cross-dataset benchmark evaluations.

8.3 Per-File-Format Marginal Value Not Yet Reported

While TURRET OS extracts metadata across five major document formats (DOCX, PDF, EXIF, EML, Git), this study evaluates the combined multi-format feature space. Isolating the precise marginal contribution of each individual document parser represents an important avenue for future ablation studies.

Table 5: ISO/IEC 27043 Forensic Readiness Attribute Coverage.

Attribute	Description	TURRET
Identification	Unique record and alert UIDs	✓
Collection	Documented harvest pipeline	✓
Acquisition	SHA-256 + BLAKE3 hash verification	✓
Preservation	Merkle chain of custody	✓
Analysis	SHAP + GNNExplainer provenance	✓
Presentation	Analyst workbench format	✓
Chain of Custody	Timestamped handler logging	✓
Integrity Verification	Ed25519 digital signature	✓
TOTAL COVERAGE		100.0%

Table 6: Computational and Storage Overhead per 1M Records.

Component	CPU (hrs)	GPU (hrs)	Storage (GB)
L1 Harvest	0.08	N/A	0.29 GB
L2 Neo4j Load	0.50	N/A	0.07 GB
L3 GNN Training	0.08	0.01	0.50 GB
L4 Evidence Packaging	0.01	N/A	0.04 GB

8.4 Single OOD Threat-Class Regime

Adversarial evaluations in this work focus on mimicry, metadata-stripping, identity-proxy spoofing, and credential-phishing evasion vectors. Emerging threat vectors, such as generative AI-driven social engineering or hardware-level DMA attacks, remain outside the current threat model scope.

8.5 Path to Industrial-Scale Tenant Evaluation

Future research will extend TURRET OS to industrial-scale cloud deployments (100,000+ endpoints), integrating distributed graph streaming engines (e.g., Apache Flink + Neo4j Fabric) to maintain sub-second detection latencies at petabyte scale.

9 Conclusion

In this paper, we introduced **TURRET OS**, a multi-layered counter-espionage framework that unifies cross-format document metadata harvest, a W3C-PROV provenance knowledge graph, a boolean rule AST engine, a spatio-temporal GNN, and Ed25519-signed Merkle evidence packaging. Evaluated across 141,725 activity records (D3 TSEC corpus) and a 90-day lab deployment pilot (D5 synthetic tenant), TURRET OS achieves an uncalibrated ROC-AUC of 0.9983, a calibrated best F1-score of 0.9747 (at threshold 0.37), an F1-score of 0.8813 at a 0.5% FPR operating point, an MCC of 0.9732, and an operational time-to-detect of < 0.1 minutes. Under adversarial red-team evaluations, TURRET OS rescues single-layer detection degradation under OOD credential-phishing attacks from 0.7363 to 0.9572 ROC-AUC. Furthermore, TURRET OS delivers 100.0% coverage across all eight ISO/IEC 27043 digital forensic readiness attributes, providing security operations centers with an automated bridge from endpoint telemetry to court-admissible forensic evidence.

Impact Statement: TURRET OS establishes a new paradigm for counter-espionage defense by demonstrating that multi-layered provenance knowledge graphs can simultaneously resolve out-of-distribution adversarial evasion and deliver court-admissible digital forensic evidence at zero operational delay.